

00utqdtmw

July 31, 2023

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: df=pd.read_csv("/content/4_drug200.csv")
df
```

```
[ ]: 
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY
..
195	56	F	LOW	HIGH	11.567	drugC
196	16	M	LOW	HIGH	12.006	drugC
197	52	M	NORMAL	HIGH	9.894	drugX
198	23	M	NORMAL	NORMAL	14.020	drugX
199	40	F	LOW	NORMAL	11.349	drugX

[200 rows x 6 columns]

```
[ ]: df.head()
```

```
[ ]: 
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY

1 DATA CLEANING AND DATA PREPROCESSING

```
[ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              200 non-null   int64
1   Sex              200 non-null   object
2   BP               200 non-null   object
3   Cholesterol      200 non-null   object
4   Na_to_K          200 non-null   float64
5   Drug             200 non-null   object
dtypes: float64(1), int64(1), object(4)
memory usage: 9.5+ KB

```

```
[ ]: df.describe()
```

```

[ ]:
      count      Age      Na_to_K
count  200.000000  200.000000
mean    44.315000   16.084485
std     16.544315    7.223956
min     15.000000    6.269000
25%     31.000000   10.445500
50%     45.000000   13.936500
75%     58.000000   19.380000
max     74.000000   38.247000

```

```
[ ]: df.columns
```

```
[ ]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

```
[ ]: df1=df.dropna(axis=1)
df1
```

```

[ ]:
      Age Sex      BP Cholesterol  Na_to_K  Drug
0     23  F    HIGH          HIGH   25.355 drugY
1     47  M    LOW           HIGH   13.093 drugC
2     47  M    LOW           HIGH   10.114 drugC
3     28  F  NORMAL          HIGH    7.798 drugX
4     61  F    LOW           HIGH   18.043 drugY
..    ..  ..    ..          ...      ...
195   56  F    LOW           HIGH   11.567 drugC
196   16  M    LOW           HIGH   12.006 drugC
197   52  M  NORMAL          HIGH    9.894 drugX
198   23  M  NORMAL        NORMAL   14.020 drugX
199   40  F    LOW        NORMAL   11.349 drugX

```

[200 rows x 6 columns]

```
[ ]: df1.columns
```

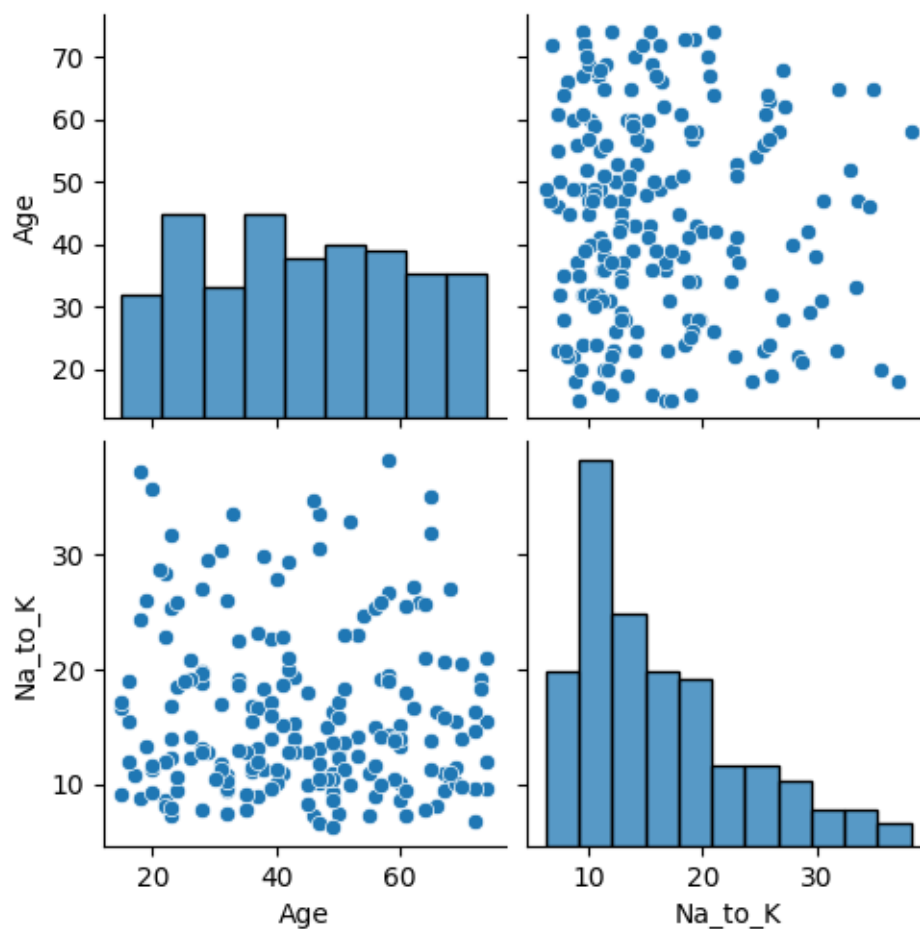
```
[ ]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

```
[ ]: df1=df1[['Age', 'Na_to_K']]
```

2 EDA AND VISUALIZATION

```
[ ]: sns.pairplot(df1)
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x7e9127466770>
```



```
[ ]: sns.distplot(df1['Na_to_K'])
```

<ipython-input-11-4b6a442fe97b>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(df1['Na_to_K'])
```

```
[ ]: <Axes: xlabel='Na_to_K', ylabel='Density'>
```



```
[ ]: sns.heatmap(df1.corr())
```

```
[ ]: <Axes: >
```



3 TO TRAIN THE MODEL AND MODEL BUILDING

```
[ ]: x=df[['Age', 'Na_to_K']]
     y=df['Na_to_K']
```

```
[ ]: from sklearn.model_selection import train_test_split
     x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)
```

```
[ ]: from sklearn.linear_model import LinearRegression
     lr=LinearRegression()
     lr.fit(x_train,y_train)
```

```
[ ]: LinearRegression()
```

```
[ ]: lr.intercept_
```


```
[ ]: -3.552713678800501e-15
```

```
[ ]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
     coeff
```

```
[ ]:      Co-efficient  
Age      1.629640e-18  
Na_to_K  1.000000e+00
```

```
[ ]: prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
[ ]: <matplotlib.collections.PathCollection at 0x7e912221cc70>
```



4 ACCURACY

```
[ ]: lr.score(x_test,y_test)
```

```
[ ]: 1.0
```

```
[ ]: lr.score(x_train,y_train)
```

```
[ ]: 1.0
```

```
[ ]: from sklearn.linear_model import Ridge,Lasso  
rr=Ridge(alpha=10)
```

```
rr.fit(x_train,y_train)
```

```
[ ]: Ridge(alpha=10)
```

```
[ ]: rr.score(x_test,y_test)
```

```
[ ]: 0.999998773006138
```

```
[ ]: rr.score(x_train,y_train)
```

```
[ ]: 0.9999988058082202
```

```
[ ]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
[ ]: Lasso(alpha=10)
```

```
[ ]: la.score(x_test,y_test)
```

```
[ ]: 0.968148097167346
```

```
[ ]: la.score(x_train,y_train)
```

```
[ ]: 0.9693674618474863
```

```
[ ]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
[ ]: ElasticNet()
```

```
[ ]: print(en.coef_)  
print(en.intercept_)
```

```
[-0.          0.98264968]  
0.28287061102032496
```

```
[ ]: prediction = en.predict(x_test)  
prediction
```

```
[ ]: array([20.82909278, 18.94437069, 15.52966305, 15.50411416, 11.10577419,  
          10.22138948, 25.62147527,  9.96098731,  9.09330765, 22.34925183,  
          14.25221847, 19.61650307, 11.42906594, 13.97511126, 10.00520655,  
           7.44147353, 17.92241502, 11.06941615, 12.9138496 , 11.64917946,  
          14.0075387 , 10.92693195, 10.99178683, 13.18211297, 16.74520071,  
          11.42611799, 10.54762917,  8.88105531, 25.80130016,  9.63179967,  
          18.01281879, 12.56107837, 18.76552845, 16.84051772,  9.79197157,
```

```
25.80621341, 18.26044651, 26.99423687, 15.17590917, 25.31587122])
```

```
[ ]: en.score(x_test,y_test)
```

```
[ ]: 0.9996869834073963
```

```
[ ]: from sklearn import metrics
print("Mean Absolute Error: ", metrics.mean_absolute_error(y_test,prediction))
print("Mean Squared Error: ", metrics.mean_squared_error(y_test,prediction))
print("Root Mean Squared Error: ", np.sqrt(metrics.
↪mean_squared_error(y_test,prediction)))
```

```
Mean Absolute Error: 0.08368614917403745
```

```
Mean Squared Error: 0.009428228438214944
```

```
Root Mean Squared Error: 0.09709906507384582
```