

evjt38y2u

July 31, 2023

1 Problem Statement

A real estate agent want help to predict the house price for regions in USA.He gave us the dataset to work on to use Linear Regression model.Create a model that helps him to estimate of what the house would sell for.

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: df=pd.read_csv("/content/10_USA_Housing.csv")
df
```

```
[ ]: 
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	\
0	79545.458574	5.682861	7.009188	
1	79248.642455	6.002900	6.730821	
2	61287.067179	5.865890	8.512727	
3	63345.240046	7.188236	5.586729	
4	59982.197226	5.040555	7.839388	
...	
4995	60567.944140	7.830362	6.137356	
4996	78491.275435	6.999135	6.576763	
4997	63390.686886	7.250591	4.805081	
4998	68001.331235	5.534388	7.130144	
4999	65510.581804	5.992305	6.792336	

	Avg. Area Number of Bedrooms	Area Population	Price	\
0	4.09	23086.800503	1.059034e+06	
1	3.09	40173.072174	1.505891e+06	
2	5.13	36882.159400	1.058988e+06	
3	3.26	34310.242831	1.260617e+06	
4	4.23	26354.109472	6.309435e+05	
...	
4995	3.46	22837.361035	1.060194e+06	
4996	4.02	25616.115489	1.482618e+06	
4997	2.13	33266.145490	1.030730e+06	
4998	5.44	42625.620156	1.198657e+06	

```
4999          4.07      46501.283803  1.298950e+06
```

```

                                Address
0  208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1  188 Johnson Views Suite 079\nLake Kathleen, CA...
2  9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3                                USS Barnett\nFPO AP 44820
4                                USNS Raymond\nFPO AE 09386
...
4995          USNS Williams\nFPO AP 30153-7653
4996          PSC 9258, Box 8489\nAPO AA 42991-3352
4997  4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998          USS Wallace\nFPO AE 73316
4999  37778 George Ridges Apt. 509\nEast Holly, NV 2...
```

```
[5000 rows x 7 columns]
```

```
[ ]: df.head()
```

```
[ ]:   Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms \
0      79545.458574          5.682861          7.009188
1      79248.642455          6.002900          6.730821
2      61287.067179          5.865890          8.512727
3      63345.240046          7.188236          5.586729
4      59982.197226          5.040555          7.839388
```

```

   Avg. Area Number of Bedrooms  Area Population      Price \
0                4.09      23086.800503  1.059034e+06
1                3.09      40173.072174  1.505891e+06
2                5.13      36882.159400  1.058988e+06
3                3.26      34310.242831  1.260617e+06
4                4.23      26354.109472  6.309435e+05
```

```

                                Address
0  208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1  188 Johnson Views Suite 079\nLake Kathleen, CA...
2  9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3                                USS Barnett\nFPO AP 44820
4                                USNS Raymond\nFPO AE 09386
```

2 Data Cleaning and Data Preprocessing

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
```

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	Avg. Area Income	5000 non-null	float64
1	Avg. Area House Age	5000 non-null	float64
2	Avg. Area Number of Rooms	5000 non-null	float64
3	Avg. Area Number of Bedrooms	5000 non-null	float64
4	Area Population	5000 non-null	float64
5	Price	5000 non-null	float64
6	Address	5000 non-null	object

dtypes: float64(6), object(1)

memory usage: 273.6+ KB

```
[ ]: df.describe()
```

```
[ ]:      Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms  \
count      5000.000000      5000.000000      5000.000000
mean      68583.108984         5.977222         6.987792
std       10657.991214         0.991456         1.005833
min       17796.631190         2.644304         3.236194
25%       61480.562388         5.322283         6.299250
50%       68804.286404         5.970429         7.002902
75%       75783.338666         6.650808         7.665871
max       107701.748378         9.519088        10.759588

      Avg. Area Number of Bedrooms  Area Population      Price
count      5000.000000      5000.000000  5.000000e+03
mean           3.981330      36163.516039  1.232073e+06
std           1.234137       9925.650114  3.531176e+05
min           2.000000       172.610686  1.593866e+04
25%           3.140000      29403.928702  9.975771e+05
50%           4.050000      36199.406689  1.232669e+06
75%           4.490000      42861.290769  1.471210e+06
max           6.500000      69621.713378  2.469066e+06
```

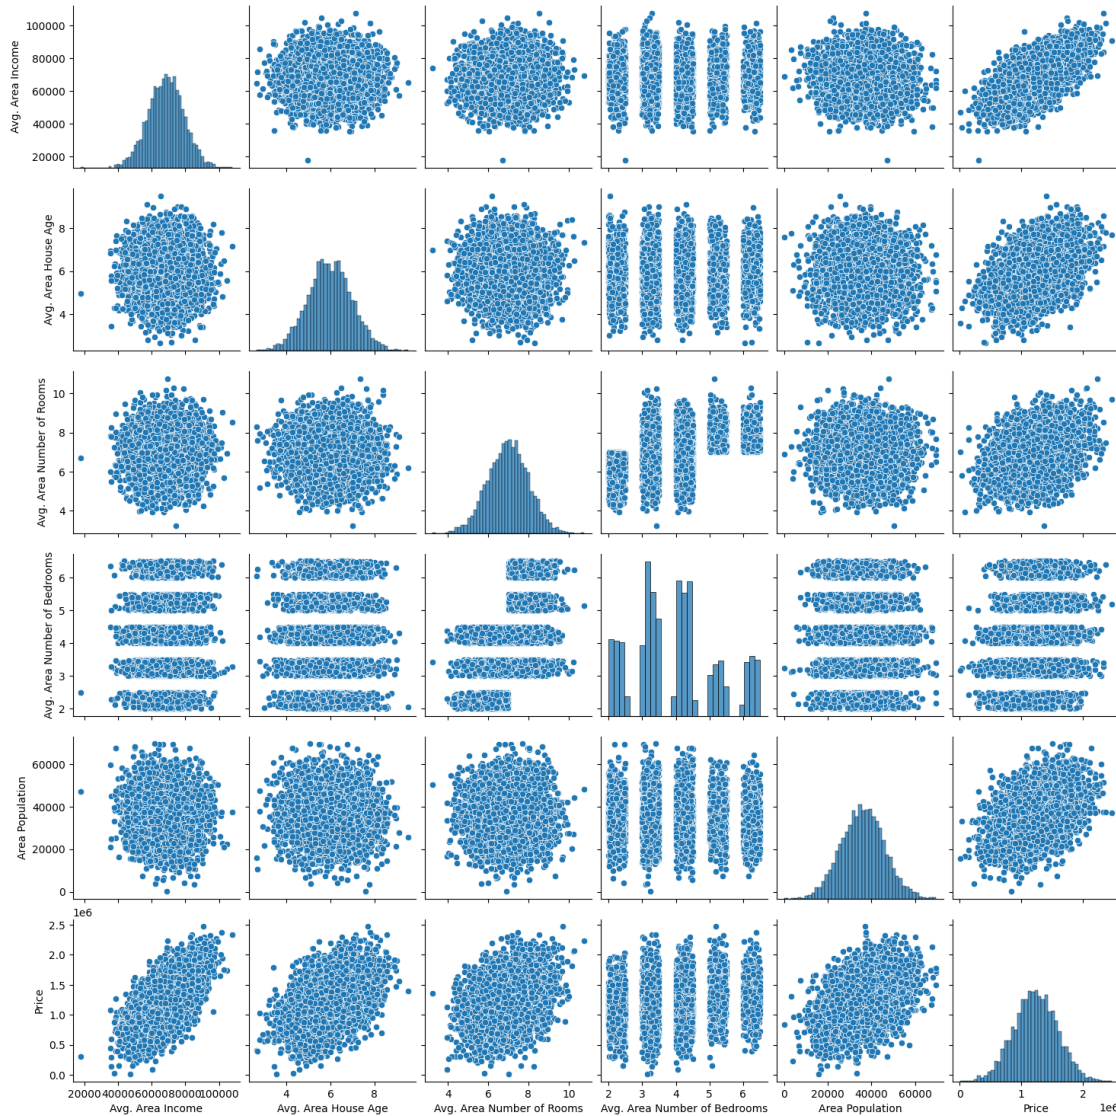
```
[ ]: df.columns
```

```
[ ]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
          'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
          dtype='object')
```

3 EDA and Visualization

```
[ ]: sns.pairplot(df)
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x7c2c00eca7d0>
```



```
[ ]: sns.distplot(df['Price'])
```

<ipython-input-8-87e11caeb2c4>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Price'])
```

```
[ ]: <Axes: xlabel='Price', ylabel='Density'>
```



```
[ ]: df1=df[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
            'Avg. Area Number of Bedrooms', 'Area Population', 'Price']]
df1
```

```
[ ]:
Avg. Area Income  Avg. Area House Age  Avg. Area Number of Rooms \
0      79545.458574      5.682861      7.009188
1      79248.642455      6.002900      6.730821
2      61287.067179      5.865890      8.512727
3      63345.240046      7.188236      5.586729
4      59982.197226      5.040555      7.839388
...
4995    60567.944140      7.830362      6.137356
4996    78491.275435      6.999135      6.576763
4997    63390.686886      7.250591      4.805081
4998    68001.331235      5.534388      7.130144
4999    65510.581804      5.992305      6.792336

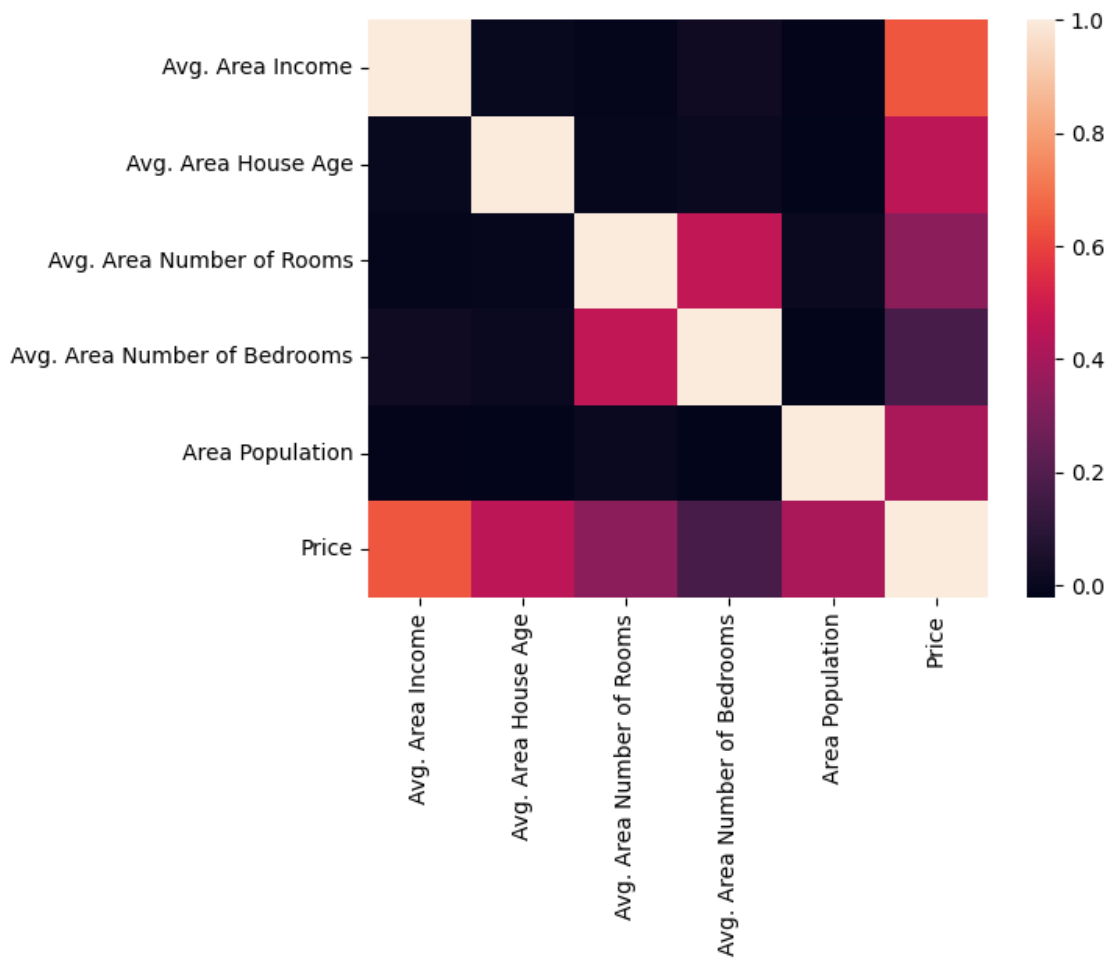
Avg. Area Number of Bedrooms  Area Population  Price
```

0	4.09	23086.800503	1.059034e+06
1	3.09	40173.072174	1.505891e+06
2	5.13	36882.159400	1.058988e+06
3	3.26	34310.242831	1.260617e+06
4	4.23	26354.109472	6.309435e+05
...
4995	3.46	22837.361035	1.060194e+06
4996	4.02	25616.115489	1.482618e+06
4997	2.13	33266.145490	1.030730e+06
4998	5.44	42625.620156	1.198657e+06
4999	4.07	46501.283803	1.298950e+06

[5000 rows x 6 columns]

```
[ ]: sns.heatmap(df1.corr())
```

```
[ ]: <Axes: >
```



To Train the Model -Model Building

We are going to train Linear Regression model; We need to split out data into two variables x and y where x is independent variable (input) and y is dependent variable on x (output) we could ignore address column as it is not required for our model

```
[ ]: x=df1[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
          'Avg. Area Number of Bedrooms', 'Area Population']]  
y=df1['Price']
```

```
[ ]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
[ ]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
[ ]: LinearRegression()
```

```
[ ]: print(lr.intercept_)
```

-2626342.537127545


```
[ ]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

```
[ ]:
```

	Co-efficient
Avg. Area Income	21.453704
Avg. Area House Age	164386.877564
Avg. Area Number of Rooms	121475.191879
Avg. Area Number of Bedrooms	1515.773131
Area Population	15.196056

```
[ ]: prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
[ ]: <matplotlib.collections.PathCollection at 0x7c2bf8444550>
```



```
[ ]: lr.score(x_test,y_test)
```

```
[ ]: 0.9211174853391793
```

```
[ ]: lr.score(x_train,y_train)
```

```
[ ]: 0.9164618662113031
```

```
[ ]: from sklearn.linear_model import Ridge,Lasso
```

```
[ ]: rr=Ridge(alpha=10)  
    rr.fit(x_train,y_train)
```

```
[ ]: Ridge(alpha=10)
```

```
[ ]: rr.score(x_test,y_test)
```

```
[ ]: 0.9210985679823257
```

```
[ ]: rr.score(x_train,y_train)
```



```
[ ]: 0.9164587767820217
```

```
[ ]: la=Lasso(alpha=10)
     la.fit(x_train,y_train)
```

```
[ ]: Lasso(alpha=10)
```

```
[ ]: la.score(x_test,y_test)
```

```
[ ]: 0.9211171631230292
```

```
[ ]: la.score(x_train,y_train)
```

```
[ ]: 0.9164618643978666
```

```
[ ]: from sklearn.linear_model import ElasticNet
     en=ElasticNet()
     en.fit(x_train,y_train)
```

```
[ ]: ElasticNet()
```

```
[ ]: en.coef_
```

```
[ ]: array([2.12263290e+01, 1.08602254e+05, 7.63297103e+04, 1.33174076e+04,
          1.51002255e+01])
```

```
[ ]: en.intercept_
```

```
[ ]: -2005738.7828657713
```

```
[ ]: prediction = en.predict(x_test)
     prediction
```

```
[ ]: array([1307199.01562841, 1064281.71599115, 1367394.6321102 , ...,
          1169293.52382855, 1386953.32730034, 1316036.19617374])
```

```
[ ]: en.score(x_test,y_test)
```

```
[ ]: 0.8840155921177741
```

```
[ ]: from sklearn import metrics
```

```
[ ]: print("Mean Absolute Error: ", metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute Error: 99675.65202139244
```

```
[ ]: print("Mean Squared Error: ", metrics.mean_squared_error(y_test,prediction))
```

Mean Squared Error: 15340780042.516378

```
[ ]: print("Root Mean Squared Error: ", np.sqrt(metrics.  
      =>mean_squared_error(y_test,prediction)))
```

Root Mean Squared Error: 123857.90262440415